



Data Mining with Weka

Logistic regression

Ian H. Witten

Logistic regression

Can do better by using prediction probabilities

Probabilities are often useful anyway ...


- ❖ Naïve Bayes produces them (obviously)
 - Open *diabetes.arff* and run *Bayes>NaiveBayes* with 90% percentage split
 - Look at columns: *actual*, *predicted*, *error*, *prob distribution*
- ❖ Other methods produce them too ...
 - Run *rules>ZeroR*. Why probabilities [0.648, 0.352] for [*tested_negative*, *tested_positive*]?
 - 90% training fold has 448 negative, 243 positive instances
 - $(448+1)/(448+1 + 243+1) = 0.648$ [cf. Laplace correction, in the Simplicity first video]
 - Run *trees>J48*
 - J48 uses probabilities internally to help with pruning

Make linear regression produce probabilities too!

Logistic regression

- ❖ Linear regression: calculate a linear function and then a threshold
- ❖ Logistic regression: estimate class probabilities directly

$$\Pr[1 \mid a_1, a_2, \dots, a_k] = 1 / (1 + \exp(-w_0 - w_1 a_1 - \dots - w_k a_k))$$



- ❖ Choose weights to maximize the log-likelihood (not minimize the squared error):

$$\sum_{i=1}^n (1 - x^{(i)}) \log(1 - \Pr[1 \mid a_1^{(1)}, a_2^{(2)}, \dots, a_k^{(k)}]) + x^{(i)} \log(\Pr[1 \mid a_1^{(1)}, a_2^{(2)}, \dots, a_k^{(k)}])$$

Logistic regression

- ❖ Open file **diabetes.arff**
- ❖ Classification-by-regression 76.8% mean of 10 runs
- ❖ cf ZeroR 65.1% 65.1%
- Naïve Bayes 76.3% 75.8%
- J48 73.8% 74.5%
- ❖ Apply **functions>Logistic** 77.2% 77.5%

- ❖ Extension to multiple classes ...
 - Perform a regression for each class?
 (like multi-response regression)
 - No. Probabilities won't sum to 1
 - Can be tackled as a joint optimization problem

Logistic regression

- ❖ Logistic regression is popular and powerful
- ❖ Uses logit transform to predict probabilities directly
 - like Naïve Bayes
- ❖ Also learned about
 - Prediction probabilities from other methods
 - How to calculate probabilities from ZeroR